

Async Phase 3

Tamilmani Manoharan
Venkatakrishnan Rajagopalan

Pseudocode

```
HashMap versionMap<(obj,attr),List<Version>>
HashMap requests<reqId, Request> # Map that stores all the original requests for restarting
HashMap recentUpdates<reqId, List<attr>> #Map store the updates that are going to be committed
HashMap cachedUpdates<obj,{attr,value}>

class Object {
    int id # SubjectId or ResourceId
    HashMap Attr<name,value>
}

class Request{
    string subId
    string resId
    string action
    string reqId
    int order
    ReqType type
    ArrayList cachedUpdates<{attr,value}>>[2] # Set of cached updates for subject and resource
    ArrayList readAttr<List<string>>>[2] # List of attribute names
    int updateIndex
    int rdOnlyIndex
    HashMap updates<attr,value> # attr name is String, value can be String or Int
}

enum ReqType {
    READ_REQ,
    WRITE_REQ,
    RESULT_RDONLY,
    RESULT_UPDATE,
    RESULT,
    UPDATEDB
}

class Message {
    ReqType type
    Request req
}

class Version {
    int rts
    int wts
}

def cachedUpdates(obj,req):
    cache = {}
    readAttr = defReadAttr(obj,req) union mightReadAttr(obj,req)
    updates = cachedUpdates[obj]
    for item in updates:
```

```

        if item.attr in readAttr:
            cache.add(item.attr,item.value)
    return cache

def latestVersionBefore(obj,attr,req.ts):
    versionList = versionMap[(obj,attr)]

    #sort version list based on timestamp
    sortedVersionList = sort(versionList)

    prev = {}
    for v in sortedVersionList:
        if(v.ts > req.ts)
            break
        prev = v

    v = prev
    if(v is empty)
        v.rts=v.wts=0
    return v

def evaluateRequest(req):
    updateIndex = -1
    updates = {}

    # check for all rules in policy files
    # change updateIndex to 1 or 2 if any one is getting updated, add changes to 'updates'

    readAttributes[1] = all attributes of 1 read so far
    readAttributes[2] = all attributes of 2 read so far

    result.decision = outcome of the policy check # True or False
    result.readAttr = readAttributes
    result.updateIndex = updateIndex
    result.updates = updates

def checkForConflicts(req):
    obj = findObject(req,req.updateIndex)

    for <attr, val> in req.updates:
        # note: if obj.attr has not been read or written in this session, then
        # v is the special version with v.rts=0 and v.wts=0.
        v = latestVersionBefore(obj,attr,req.ts)
        if v.rts > req.ts:
            return true
    return false

def restart(req):
    originalReq = requests[req.reqId]
    obj = findObject(originalReq,1)
    coordinatorId = findCoordinator(obj)
    req.order=1
    sendRequest(originalReq,originalReq.type,coordinatorId)

Client() {
    Message msg
    while(msg = recvMsg())
    {
        if(msg.type == READ_REQ or WRITE_REQ)
        {
            req = msg.req
            req.type = msg.type

```

```

    obj = findObject(req,1)
    coordinatorId = findCoordinator(obj)
    req.order=1
    sendRequest(req,msg.type,coordinatorId)
}
else if(msg.type == RESULT)
    sendResultToApp(result)
}

Coordinator() {
    Message msg
    while( msg = recvMsg())
    {
        if(msg.type == READ_REQ or WRITE_REQ)
        {
            req = msg.req
            requests[req.reqId] = req # Store original request, for retrieving and restarting later
            order = req.order
            obj = findObject(req,order)

            if(order == 1) {

                # To prevent starvation of write requests
                readAttr = defReadAttr(obj,req) union mightReadAttr(obj,req)
                for item in recentUpdates:
                    await(no attr in item.attr is present in readAttr)

                req.ts = now()
            }

            if(msg.type == READ_REQ)
            {
                for attr in defReadAttr(obj,req):
                    latestVersionBefore(obj,attr,req.ts).rts = req.ts

                for attr in mightReadAttr(obj,req):
                    latestVersionBefore(obj,attr,req.ts).pendingMightRead.add(req.id)
            }
            else if(msg.type==WRITE_REQ)
            {
                for attr in defReadAttr(obj,req) union mightReadAttr(obj,req)
                    v = latestVersionBefore(obj,attr,req.ts)
                    v.pendingMightRead.add(req.id)
            }

            req.cachedUpdates[order] = cachedUpdates(obj,req)

            if(order==1)
                obj = findObject(req,2)
                coordinatorId = findCoordinator(obj)
                req.order=2
                send(req,msg.type,coordinatorId)
            else
                workerId = getWorker(obj)
                send(req,msg.type,workerId)
        }
        else if(msg.type == RESULT_RDONLY)
        {
            req = msg.req
            obj = findObject(req,req.order)
            for attr in mightReadAttr(obj,req)

```

```

        v = latestVersionBefore(obj,attr,req.ts)
        v.pendingMightRead.remove(req.id)
        if attr in req.readAttr[i]:
            v.rts = req.ts
    }
else if(msg.type == RESULT_UPDATE)
{
    obj = findObject(req,req.updateIndex)

    conflict = checkForConflicts(req)

    # Store recent updates to check for starvation in incoming new read requests
    forall <attr,val> in req.updates:
        recentUpdates[req.id].append(attr)

    if not conflict:

        # wait for relevant pending reads to complete
        await (forall <attr,val> in req.updates:
            latestVersionBefore(x,attr,req.ts).pendingMightRead is empty
            or contains only an entry for req)
        # check again for conflicts
        conflict = checkForConflicts(req)
        if not conflict:
            # commit the updates
            send(req,UPDATEDB,dbID)

            #create new version and append to version map
            forall <attr,val> in req.updates:
                cachedUpdates[obj].append(<attr,val>)
                Version newVersion = Version()
                newVersion.rts = 0
                newVersion.wts = req.ts
                versionMap[(obj,attr)].append(newVersion)

            # update read timestamps

            for attr in defReadAttr(x,req) union mightReadAttr(x,req):
                v = latestVersionBefore(x,attr,req.ts)
                v.pendingMightRead.remove(req.id)
                if attr in req.readAttr[req.updatedObj]:
                    v.rts = req.ts

            #clearing the recent updates after committing
            recentUpdates[req.id] = None

            send(req,RESULT,req.clientId)

            obj = findObject(req,req.rdOnlyIndex)
            coordinatorId = findCoordinator(obj)
            send(req, RESULT_RDONLY, coordinatorId)
        else:
            restart(req)
    else:
        restart(req)
}
}
}

Worker() {
    Message msg
    while(msg=recvMsg())

```

```

{
    req = msg.req
    result = evaluateRequest(req)
    req.decision = result.decision

    for i = 1 to 2:
        req.readAttr[i] = result.readAttr[findObject(req,i)]

    #For Read only requests, result.updates will be empty and result.updateIndex will be -1
    req.updateIndex = result.updateIndex
    req.updates = result.updates

    if(req.updateIndex == -1)
        send(req,RESULT,req.clientId)

        for i = 1 to 2:
            obj = findObject(req,i)
            req.order = i
            coordinatorId = findCoordinator(obj)
            send(req,RESULT_RDONLY,coordinatorId)
        else
            obj = findObject(req,req.updateIndex)
            coordinatorId = findCoordinator(obj)
            req.rdOnlyIndex = 3 - updateIndex # since only 1 or 2 is possible
            send(req,RESULT_UPDATE,coordinatorId)
    }
}

```
